

The Apple PSI System

[Bhowmick et al., 2021]

Alessandro Baccharini

University at Buffalo

anbaccar@buffalo.edu

December 15, 2021

Table of Contents

1 Motivations

2 Protocol Description

3 Conclusions

Motivations

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
 - Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
 - Aimed to be packaged with iOS 15 and iPadOS 15.
 - Very poorly received in media and tech communities.
-
- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
 - Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
 - Aimed to be packaged with iOS 15 and iPadOS 15.
 - Very poorly received in media and tech communities.
-
- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
 - Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
 - Aimed to be packaged with iOS 15 and iPadOS 15.
 - Very poorly received in media and tech communities.
-
- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
 - Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
 - Aimed to be packaged with iOS 15 and iPadOS 15.
 - Very poorly received in media and tech communities.
-
- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
 - Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
 - Aimed to be packaged with iOS 15 and iPadOS 15.
 - Very poorly received in media and tech communities.
-

Why Apple's child safety updates are so controversial

Apple is trying to balance child safety and privacy, but some experts say the company is going too far.

- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
- Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
- Aimed to be packaged with iOS 15 and iPadOS 15.
- Very poorly received in media and tech communities.

Aug 24, 2021, 09:00am EDT | 404,821 views

Researchers Label Apple's CSAM Detection System 'Dangerous'

- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
- Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
- Aimed to be packaged with iOS 15 and iPadOS 15.
- Very poorly received in media and tech communities.

Apple's CSAM detection tech is under fire — again



Zack Whittaker @zackwhittaker / 2:28 PM EDT • August 18, 2021

Comment

- September 2021 – Apple postpones rollout indefinitely.

Why?

- August 2021 – Apple unveils plans for new Child Sexual Abuse Material (CSAM) detection system.
- Designed to automatically detect known CSAM images stored in iCloud, and report the users to authorities.
- Aimed to be packaged with iOS 15 and iPadOS 15.
- Very poorly received in media and tech communities.

Apple's CSAM detection tech is under fire — again

Zack Whittaker @zackwhittaker / 2:28 PM EDT • August 18, 2021

 Comment

- September 2021 – Apple postpones rollout indefinitely.

What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.

What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.

What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.

What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.

What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.


What security goals do “we” want?

- Server cannot recover the user’s matched photos without exceeding some threshold.
- False positives are impossible.
- No information is learned about non-matched images.
- User cannot learn any information from the CSAM database.
- User cannot identify which images were flagged as CSAM by the system.

- Different from our standard notion of hash functions.
- Insensitive to small perturbations (cropping, rotation, mirroring, watermarking).

NeuralHash

- Different from our standard notion of hash functions.
- Insensitive to small perturbations (cropping, rotation, mirroring, watermarking).



[Struppek et al., 2021]

- Contains some collision-related issues [Athalye, 2021]...

- Contains some collision-related issues [Athalye, 2021]...



- Contains some collision-related issues [Athalye, 2021]...



```
$ python nnhash.py cat.png  
59a34eabe31910abfb06f308  
$ python nnhash.py dog.png  
59a34eabe31910abfb06f308
```

A Crash Course in Private Set Intersection (PSI)

- Let \mathcal{U} be the universe of all possible image hashes.
- $X \subseteq \mathcal{U}$ is set of image hashes we want to match against, stored on the server.
- A client has a list of m triples

$$\bar{Y} = ((y_1, id_1, ad_1), \dots, (y_m, id_m, ad_m)) \in (\mathcal{U} \times \mathcal{ID} \times \mathcal{D})^m,$$

where $y \in \mathcal{U}$ is the hash of an image, a unique identifier $id \in \mathcal{ID}$, and some associated data $ad \in \mathcal{D}$.

- When the protocol terminates, the server learns the identifiers and associated data of the intersection of \bar{Y} and X , namely $id(\bar{Y} \cap X)$

A Crash Course in Private Set Intersection (PSI)

- Let \mathcal{U} be the universe of all possible image hashes.
- $X \subseteq \mathcal{U}$ is set of image hashes we want to match against, stored on the server.
- A client has a list of m triples

$$\bar{Y} = ((y_1, id_1, ad_1), \dots, (y_m, id_m, ad_m)) \in (\mathcal{U} \times \mathcal{ID} \times \mathcal{D})^m,$$

where $y \in \mathcal{U}$ is the hash of an image, a unique identifier $id \in \mathcal{ID}$, and some associated data $ad \in \mathcal{D}$.

- When the protocol terminates, the server learns the identifiers and associated data of the intersection of \bar{Y} and X , namely $id(\bar{Y} \cap X)$

A Crash Course in Private Set Intersection (PSI)

- Let \mathcal{U} be the universe of all possible image hashes.
- $X \subseteq \mathcal{U}$ is set of image hashes we want to match against, stored on the server.
- A client has a list of m triples

$$\bar{Y} = ((y_1, id_1, ad_1), \dots, (y_m, id_m, ad_m)) \in (\mathcal{U} \times \mathcal{ID} \times \mathcal{D})^m,$$

where $y \in \mathcal{U}$ is the hash of an image, a unique identifier $id \in \mathcal{ID}$, and some associated data $ad \in \mathcal{D}$.

- When the protocol terminates, the server learns the identifiers and associated data of the intersection of \bar{Y} and X , namely $id(\bar{Y} \cap X)$

A Crash Course in Private Set Intersection (PSI)

- Let \mathcal{U} be the universe of all possible image hashes.
- $X \subseteq \mathcal{U}$ is set of image hashes we want to match against, stored on the server.
- A client has a list of m triples

$$\bar{Y} = ((y_1, id_1, ad_1), \dots, (y_m, id_m, ad_m)) \in (\mathcal{U} \times \mathcal{ID} \times \mathcal{D})^m,$$

where $y \in \mathcal{U}$ is the hash of an image, a unique identifier $id \in \mathcal{ID}$, and some associated data $ad \in \mathcal{D}$.

- When the protocol terminates, the server learns the identifiers and associated data of the intersection of \bar{Y} and X , namely $id(\bar{Y} \cap X)$

Two PSI Protocols

Threshold PSI-AD

Add a threshold parameter t , such that if $|id(\bar{Y} \cap X)| \leq t$, the server learns only the id 's. If $|id(\bar{Y} \cap X)| > t$, then the server learns the associated data for all identifiers in the intersection.

Two PSI Protocols

Threshold PSI-AD

Add a threshold parameter t , such that if $|id(\bar{Y} \cap X)| \leq t$, the server learns only the id 's. If $|id(\bar{Y} \cap X)| > t$, then the server learns the associated data for all identifiers in the intersection.

Fuzzy Threshold PSI-AD

Extension of prior scheme, but adds “synthetic matches” so the server does not know the number of matches in the intersection before the threshold t is exceeded.

Two PSI Protocols

Threshold PSI-AD

Add a threshold parameter t , such that if $|id(\bar{Y} \cap X)| \leq t$, the server learns only the id 's. If $|id(\bar{Y} \cap X)| > t$, then the server learns the associated data for all identifiers in the intersection.

Fuzzy Threshold PSI-AD

Extension of prior scheme, but adds “synthetic matches” so the server does not know the number of matches in the intersection before the threshold t is exceeded.

Protocol Description

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

Server Setup

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

1. Remove any duplicates from X , and let $n = |X|$.
2. Construct a hash table T :
 - Let $n' \geq n$ be the size of the table (minimize collisions).
 - Choose hash function $h : \mathcal{U} \rightarrow \{1, \dots, n'\}$ (SHA256 modulo n').
 - Insert elements of X into T , each cell should have at most one element.
3. Choose a random nonzero $\alpha \in \mathbb{F}_q$, compute $L = G^\alpha \in \mathbb{G}$, where \mathbb{G} is a DH group modulo prime p (2048-bit) with a fixed generator $G = 2$.
4. For $i = 1$ to n' do:
 - If $T[i]$ is non-empty, set $P_i = H(T[i])^\alpha \in \mathbb{G}$, where $T[i] \in X \subseteq \mathcal{U}$, and $H : \mathcal{U} \rightarrow \mathbb{G}$ (SHA256 modulo p).
 - If $T[i]$ is empty, choose a random $P_i \in \mathbb{G}$.
5. set $pdata = (L, P_1, \dots, P_{n'})$.

1. Obtain *pdata* from the server.
2. Generate keys:
 - $adkey \leftarrow \mathcal{K}'$ for encryption scheme (Enc, Dec) .
 - We use AES128-GCM for its “random key robustness” property.
 - $\text{Dec}(\text{Enc}(k, m), k')$ should fail, where $k \neq k'$ are independent random keys.
 - $fkey \leftarrow \mathcal{K}''$ for the PRF $F : \mathcal{K}'' \times \mathcal{ID} \rightarrow \mathbb{F}_{\text{Sh}}$.
 - Initialize threshold *Shamir secret sharing* for *adkey*:

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_tx^t,$$

where $a_0 = adkey$ is the secret. Reconstruction involves Lagrange interpolation.

Client Setup

1. Obtain *pdata* from the server.
2. Generate keys:
 - $adkey \leftarrow \mathcal{K}'$ for encryption scheme (Enc, Dec) .
 - We use AES128-GCM for its “random key robustness” property.
 - $\text{Dec}(\text{Enc}(k, m), k')$ should fail, where $k \neq k'$ are independent random keys.
 - $fkey \leftarrow \mathcal{K}''$ for the PRF $F : \mathcal{K}'' \times \mathcal{ID} \rightarrow \mathbb{F}_{\text{Sh}}$.
 - Initialize threshold *Shamir secret sharing* for *adkey*:

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_tx^t,$$

where $a_0 = adkey$ is the secret. Reconstruction involves Lagrange interpolation.

Client Setup

1. Obtain *pdata* from the server.
2. Generate keys:
 - $adkey \leftarrow \mathcal{K}'$ for encryption scheme (Enc, Dec) .
 - We use AES128-GCM for its “random key robustness” property.
 - $\text{Dec}(\text{Enc}(k, m), k')$ should fail, where $k \neq k'$ are independent random keys.
 - $fkey \leftarrow \mathcal{K}''$ for the PRF $F : \mathcal{K}'' \times \mathcal{ID} \rightarrow \mathbb{F}_{\text{Sh}}$.
 - Initialize threshold *Shamir secret sharing* for *adkey*:

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_tx^t,$$

where $a_0 = adkey$ is the secret. Reconstruction involves Lagrange interpolation.

Client Setup

1. Obtain *pdata* from the server.
2. Generate keys:
 - $adkey \leftarrow \mathcal{K}'$ for encryption scheme (Enc, Dec) .
 - We use AES128-GCM for its “random key robustness” property.
 - $\text{Dec}(\text{Enc}(k, m), k')$ should fail, where $k \neq k'$ are independent random keys.
 - $fkey \leftarrow \mathcal{K}''$ for the PRF $F : \mathcal{K}'' \times \mathcal{ID} \rightarrow \mathbb{F}_{\text{Sh}}$.
 - Initialize threshold *Shamir secret sharing* for *adkey*:

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_tx^t,$$

where $a_0 = adkey$ is the secret. Reconstruction involves Lagrange interpolation.

1. Obtain *pdata* from the server.
2. Generate keys:
 - $adkey \leftarrow \mathcal{K}'$ for encryption scheme (Enc, Dec) .
 - We use AES128-GCM for its “random key robustness” property.
 - $\text{Dec}(\text{Enc}(k, m), k')$ should fail, where $k \neq k'$ are independent random keys.
 - $fkey \leftarrow \mathcal{K}''$ for the PRF $F : \mathcal{K}'' \times \mathcal{ID} \rightarrow \mathbb{F}_{\text{Sh}}$.
 - Initialize threshold *Shamir secret sharing* for *adkey*:

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_tx^t,$$

where $a_0 = adkey$ is the secret. Reconstruction involves Lagrange interpolation.

Client Voucher Generation on Input Triple (y, id, ad)

1. Encrypt ad as $adct \leftarrow \text{Enc}(adkey, ad)$, and all $adct$ must be the same length.
2. Compute $x = F(fkey, id) \in \mathbb{F}_{Sh}$.
3. Generate a share $sh = (x, f(x)) \in \mathbb{F}_{Sh}$ of $adkey$ (guarantees duplicate triples with the same id will produce the same sh).
4. Choose a random key $rkey \leftarrow \mathcal{K}'$ and compute $rct \leftarrow \text{Enc}(rkey, (adct, sh))$.

Client Voucher Generation on Input Triple (y, id, ad)

1. Encrypt ad as $adct \leftarrow \text{Enc}(adkey, ad)$, and all $adct$ must be the same length.
2. Compute $x = F(fkey, id) \in \mathbb{F}_{Sh}$.
3. Generate a share $sh = (x, f(x)) \in \mathbb{F}_{Sh}$ of $adkey$ (guarantees duplicate triples with the same id will produce the same sh).
4. Choose a random key $rkey \leftarrow \mathcal{K}'$ and compute $rct \leftarrow \text{Enc}(rkey, (adct, sh))$.

Client Voucher Generation on Input Triple (y, id, ad)

1. Encrypt ad as $adct \leftarrow \text{Enc}(adkey, ad)$, and all $adct$ must be the same length.
2. Compute $x = F(fkey, id) \in \mathbb{F}_{Sh}$.
3. Generate a share $sh = (x, f(x)) \in \mathbb{F}_{Sh}$ of $adkey$ (guarantees duplicate triples with the same id will produce the same sh).
4. Choose a random key $rkey \leftarrow \mathcal{K}'$ and compute $rct \leftarrow \text{Enc}(rkey, (adct, sh))$.

Client Voucher Generation on Input Triple (y, id, ad)

1. Encrypt ad as $adct \leftarrow \text{Enc}(adkey, ad)$, and all $adct$ must be the same length.
2. Compute $x = F(fkey, id) \in \mathbb{F}_{Sh}$.
3. Generate a share $sh = (x, f(x)) \in \mathbb{F}_{Sh}$ of $adkey$ (guarantees duplicate triples with the same id will produce the same sh).
4. Choose a random key $rkey \leftarrow \mathcal{K}'$ and compute $rct \leftarrow \text{Enc}(rkey, (adct, sh))$.

Client Voucher Generation on Input Triple (y, id, ad)

5. Compute $w = h(y) \in \{1, \dots, n'\}$.
6. Sample random $\beta, \gamma \in \mathbb{F}_q$, and use P_w, L from $pdata$ to compute:

$$Q = H(y)^\beta \cdot G^\gamma \text{ and } S = P_w^\beta \cdot L^\gamma,$$

where if $y = T[w]$, then $P_w = H(y)^\alpha$ and $S = Q^\alpha$ (DH random self reduction).

7. Compute $ct \leftarrow \text{Enc}(H'(S), rkey)$, where $H' : \mathbb{G} \rightarrow \mathcal{K}'$ (HKDF with SHA256).
8. Send $voucher = (id, Q, ct, rct)$ to the server.

Client Voucher Generation on Input Triple (y, id, ad)

5. Compute $w = h(y) \in \{1, \dots, n'\}$.
6. Sample random $\beta, \gamma \in \mathbb{F}_q$, and use P_w, L from $pdata$ to compute:

$$Q = H(y)^\beta \cdot G^\gamma \text{ and } S = P_w^\beta \cdot L^\gamma,$$

where if $y = T[w]$, then $P_w = H(y)^\alpha$ and $S = Q^\alpha$ (DH random self reduction).

7. Compute $ct \leftarrow \text{Enc}(H'(S), rkey)$, where $H' : \mathbb{G} \rightarrow \mathcal{K}'$ (HKDF with SHA256).
8. Send $voucher = (id, Q, ct, rct)$ to the server.

Client Voucher Generation on Input Triple (y, id, ad)

5. Compute $w = h(y) \in \{1, \dots, n'\}$.
6. Sample random $\beta, \gamma \in \mathbb{F}_q$, and use P_w, L from $pdata$ to compute:

$$Q = H(y)^\beta \cdot G^\gamma \text{ and } S = P_w^\beta \cdot L^\gamma,$$

where if $y = T[w]$, then $P_w = H(y)^\alpha$ and $S = Q^\alpha$ (DH random self reduction).

7. Compute $ct \leftarrow \text{Enc}(H'(S), rkey)$, where $H' : \mathbb{G} \rightarrow \mathcal{K}'$ (HKDF with SHA256).
8. Send $voucher = (id, Q, ct, rct)$ to the server.

Client Voucher Generation on Input Triple (y, id, ad)

5. Compute $w = h(y) \in \{1, \dots, n'\}$.
6. Sample random $\beta, \gamma \in \mathbb{F}_q$, and use P_w, L from $pdata$ to compute:

$$Q = H(y)^\beta \cdot G^\gamma \text{ and } S = P_w^\beta \cdot L^\gamma,$$

where if $y = T[w]$, then $P_w = H(y)^\alpha$ and $S = Q^\alpha$ (DH random self reduction).

7. Compute $ct \leftarrow \text{Enc}(H'(S), rkey)$, where $H' : \mathbb{G} \rightarrow \mathcal{K}'$ (HKDF with SHA256).
8. Send $voucher = (id, Q, ct, rct)$ to the server.

Server Voucher Processing

1. Initialize empty set *SHARES* and an empty list *IDLIST*.
2. For each voucher (id, Q, ct, rct) received, do:
 - Append *id* to *IDLIST*.
 - Compute $\hat{S} = Q^\alpha \in \mathbb{G}$,
 - Set $rkey = \text{Dec}(H'(\hat{S}), ct)$.
 - Set $(adct, sh) = \text{Dec}(rkey, rct)$.
 - If either decryptions “fails”, *y* is a non-match, and ignore the voucher.
 - Otherwise, we found a match and add $(id, adct, sh)$ to *SHARES*.

Server Voucher Processing

1. Initialize empty set $SHARES$ and an empty list $IDLIST$.
2. For each voucher (id, Q, ct, rct) received, do:
 - Append id to $IDLIST$.
 - Compute $\hat{S} = Q^\alpha \in \mathbb{G}$,
 - Set $rkey = \text{Dec}(H'(\hat{S}), ct)$.
 - Set $(adct, sh) = \text{Dec}(rkey, rct)$.
 - If either decryptions “fails”, y is a non-match, and ignore the voucher.
 - Otherwise, we found a match and add $(id, adct, sh)$ to $SHARES$.

Server Voucher Processing

1. Initialize empty set $SHARES$ and an empty list $IDLIST$.
2. For each voucher (id, Q, ct, rct) received, do:
 - Append id to $IDLIST$.
 - Compute $\hat{S} = Q^\alpha \in \mathbb{G}$,
 - Set $rkey = \text{Dec}(H'(\hat{S}), ct)$.
 - Set $(adct, sh) = \text{Dec}(rkey, rct)$.
 - If either decryptions “fails”, y is a non-match, and ignore the voucher.
 - Otherwise, we found a match and add $(id, adct, sh)$ to $SHARES$.

Server Voucher Processing

1. Initialize empty set $SHARES$ and an empty list $IDLIST$.
2. For each voucher (id, Q, ct, rct) received, do:
 - Append id to $IDLIST$.
 - Compute $\hat{S} = Q^\alpha \in \mathbb{G}$,
 - Set $rkey = \text{Dec}(H'(\hat{S}), ct)$.
 - Set $(adct, sh) = \text{Dec}(rkey, rct)$.
 - If either decryptions “fails”, y is a non-match, and ignore the voucher.
 - Otherwise, we found a match and add $(id, adct, sh)$ to $SHARES$.

3. Let t' denote the number of *unique* shares in $SHARES$, and t' should equal the size of $id(\bar{Y} \cap X)$.
 - If $t' \leq t$, let $OUTSET$ be the set of identifiers in $SHARES$.
 - If $t' > t$, do:
 - Use $(t + 1)$ shares to reconstruct $adkey \in \mathcal{K}'$.
 - Initialize $OUTSET = \{\emptyset\}$.
 - For each triple $(id, adct, sh) \in SHARES$, compute $ad = Dec(adkey, adct)$. If it fails, discard the voucher. Otherwise, add (id, ad) to $OUTLIST$.
 - Output $IDLIST$ and $OUTSET$.

3. Let t' denote the number of *unique* shares in $SHARES$, and t' should equal the size of $id(\bar{Y} \cap X)$.
 - If $t' \leq t$, let $OUTSET$ be the set of identifiers in $SHARES$.
 - If $t' > t$, do:
 - Use $(t + 1)$ shares to reconstruct $adkey \in \mathcal{K}'$.
 - Initialize $OUTSET = \{\emptyset\}$.
 - For each triple $(id, adct, sh) \in SHARES$, compute $ad = Dec(adkey, adct)$. If it fails, discard the voucher. Otherwise, add (id, ad) to $OUTLIST$.
 - Output $IDLIST$ and $OUTSET$.

3. Let t' denote the number of *unique* shares in $SHARES$, and t' should equal the size of $id(\bar{Y} \cap X)$.
 - If $t' \leq t$, let $OUTSET$ be the set of identifiers in $SHARES$.
 - If $t' > t$, do:
 - Use $(t + 1)$ shares to reconstruct $adkey \in \mathcal{K}'$.
 - Initialize $OUTSET = \{\emptyset\}$.
 - For each triple $(id, adct, sh) \in SHARES$, compute $ad = \text{Dec}(adkey, adct)$. If it fails, discard the voucher. Otherwise, add (id, ad) to $OUTLIST$.
 - Output $IDLIST$ and $OUTSET$.

3. Let t' denote the number of *unique* shares in $SHARES$, and t' should equal the size of $id(\bar{Y} \cap X)$.
 - If $t' \leq t$, let $OUTSET$ be the set of identifiers in $SHARES$.
 - If $t' > t$, do:
 - Use $(t + 1)$ shares to reconstruct $adkey \in \mathcal{K}'$.
 - Initialize $OUTSET = \{\emptyset\}$.
 - For each triple $(id, adct, sh) \in SHARES$, compute $ad = \text{Dec}(adkey, adct)$. If it fails, discard the voucher. Otherwise, add (id, ad) to $OUTLIST$.
 - Output $IDLIST$ and $OUTSET$.

(Brief) Discussion

- Protocol is correct if the client and server adhere to the protocol (proof omitted for obvious reasons).
- Using “simpler” constructions guarantees the same level of security as the original protocol (potentially for the price of degraded performance).
- Construction naturally extends to ftPSI-AD, requires novel primitives.
 - Detectable hash functions, hashing to elliptic curves, etc.

(Brief) Discussion

- Protocol is correct if the client and server adhere to the protocol (proof omitted for obvious reasons).
- Using “simpler” constructions guarantees the same level of security as the original protocol (potentially for the price of degraded performance).
- Construction naturally extends to ftPSI-AD, requires novel primitives.
 - Detectable hash functions, hashing to elliptic curves, etc.

(Brief) Discussion

- Protocol is correct if the client and server adhere to the protocol (proof omitted for obvious reasons).
- Using “simpler” constructions guarantees the same level of security as the original protocol (potentially for the price of degraded performance).
- Construction naturally extends to ftPSI-AD, requires novel primitives.
 - Detectable hash functions, hashing to elliptic curves, etc.

(Brief) Discussion

- Protocol is correct if the client and server adhere to the protocol (proof omitted for obvious reasons).
- Using “simpler” constructions guarantees the same level of security as the original protocol (potentially for the price of degraded performance).
- Construction naturally extends to ftPSI-AD, requires novel primitives.
 - Detectable hash functions, hashing to elliptic curves, etc.

Conclusions

Conclusions

- Presented Apple's PSI system for CSAM detection.
- The protocol is cryptographically sound, and meets the security goals specified earlier.

Conclusions

- Presented Apple's PSI system for CSAM detection.
- The protocol is cryptographically sound, and meets the security goals specified earlier.

Conclusions

- So why is something like this still bad?
- What are the implications of this system?

Conclusions

- So why is something like this still bad?
- What are the implications of this system?

References



Athalye, A. (2021).

NeuralHash Collider.

<https://github.com/anishathalye/neural-hash-collider>.

original-date: 2021-08-19T00:06:20Z.



Bhowmick, A., Boneh, D., Myers, S., and Tarpe, K. T. K. (2021).

The Apple PSI System.

<https://www.apple.com/child-safety/pdf/>

[Apple_PSI_System_Security_Protocol_and_Analysis.pdf](#).



Struppek, L., Hintersdorf, D., Neider, D., and Kersting, K. (2021).

Learning to break deep perceptual hashing: The use case neuralhash.

arXiv preprint arXiv:2111.06628.

Thank you!

Questions?